

# Text Classification


Lena Voita

---

Lecture-blog and lots of additional materials are here:  
[https://lena-voita.github.io/nlp\\_course/text\\_classification.html](https://lena-voita.github.io/nlp_course/text_classification.html)


NLP Course For You 

# What is going to happen:

- Examples of classification tasks
  - General View: Features + Classifier
  - Models: Generative vs Discriminative
  - Classical Methods
  - Neural Methods
  - Multi-Label Classification
  - Practical Tips
  -  Analysis and Interpretability
- 
- High-Level View
  - Training: Cross-Entropy
  - Models: (Weighted) BOW
  - Models: Convolutional
  - Models: Recurrent

# Basics: Recurrent Neural Networks


- recurrent cell




For more details of RNN basics, look at the [Colah's blog post](#).

# Basics: Recurrent Neural Networks

- recurrent cell



- vanilla RNN




For more details of RNN basics, look at the [Colah's blog post](#).

# Recurrent Models for Text Classification


We need a model that can produce a **fixed-sized** vector for inputs of **different lengths**.

- simple: read a text, take the final state




# The Problems with Recurrence

- Vanishing and Exploding Gradients
  - Lots of matrix multiplications either blow up or diminish the activity going through the network




# The Problems with Recurrence

- Vanishing and Exploding Gradients
  - Lots of matrix multiplications either blow up or diminish the activity going through the network
- Credit assignment
  - While it is possible for the network to be sensitive to when a token is seen it is very subtle and position information is often lost in the final vector
  - There is also a recency bias – recent information dominates the combined vector representation while earlier information is forgotten




# The Problems with Recurrence

- Vanishing and Exploding Gradients
  - Lots of matrix multiplications either blow up or diminish the activity going through the network
- Credit assignment
  - While it is possible for the network to be sensitive to when a token is seen it is very subtle and position information is often lost in the final vector
  - There is also a recency bias – recent information dominates the combined vector representation while earlier information is forgotten
- Fixed vector size:
  - The size of the combined vector does not adjust to the length of the sentence




# Recurrent Models for Text Classification

- Multiple layers: feed the states from one RNN to the next




# Recurrent Models for Text Classification


- bidirectional: use final states from forward and backward RNNs




# What is going to happen:

- Examples of classification tasks
  - General View: Features + Classifier
  - Models: Generative vs Discriminative
  - Classical Methods
  - Neural Methods
  - Multi-Label Classification
  - Practical Tips
  -  Analysis and Interpretability
- 
- High-Level View
  - Training: Cross-Entropy
  - Models: (Weighted) BOW
  - Models: Convolutional
  - Models: Recurrent

# What is going to happen:

- Examples of classification tasks
  - General View: Features + Classifier
  - Models: Generative vs Discriminative
  - Classical Methods
  - Neural Methods
  - Multi-Label Classification
  - Practical Tips
  -  Analysis and Interpretability
- 
- Word Embeddings
  - Data Augmentation

# Word Embeddings: How to Deal with them?




Input word embeddings:


- Train from scratch
- Take pretrained (Word2Vec, GloVe)
- Initialize with pretrained, then fine-tune

# Which data do we have?

Training data for text classification (labeled)



Training data for word embeddings (unlabeled)




- Not huge, or not diverse, or both
- Domain: task-specific

- Huge diverse corpus (e.g., Wikipedia)
- Domain: general

# Word Embeddings: How to Deal with them?

- Train from scratch


What they will know:



May be not enough  
to learn relationships  
between words

- Take pretrained  
(Word2Vec, GloVe)


What they will know:



Know relationships between words,  
but are **not specific to the task**

- Initialize with pretrained,  
then fine-tune

What they will know:




Know relationships between  
words and adapted for the task

# Word Embeddings: How to Deal with them?

- Train from scratch


What they will know:



May be not enough  
to learn relationships  
between words

- Take pretrained  
(Word2Vec, GloVe)


What they will know:



Know relationships between words,  
but are **not specific to the task**

- Initialize with pretrained,  
then fine-tune

What they will know:




Know relationships between  
words and adapted for the task


“Transfer” knowledge from a huge unlabeled  
corpus to your task-specific model

We’ll learn more about this later in the course!

# What is going to happen:

- Examples of classification tasks
  - General View: Features + Classifier
  - Models: Generative vs Discriminative
  - Classical Methods
  - Neural Methods
  - Multi-Label Classification
  - Practical Tips
  -  Analysis and Interpretability
- 
- Word Embeddings
  - Data Augmentation

# What is going to happen:

- Examples of classification tasks
- General View: Features + Classifier
- Models: Generative vs Discriminative
- Classical Methods
- Neural Methods
- Multi-Label Classification
- Practical Tips
-  Analysis and Interpretability



- Word Embeddings
- Data Augmentation

# Data Augmentation: Get More Data for Free


Data augmentation alters your dataset in different ways to get alternative versions of the same training example.

It can increase:

- the amount of data
- diversity of data

# Data Augmentation for Images

- Flipping an image
- Geometrical transformations (rotation, stretching, zoom in/out)
- Covering some parts with patches



# Data Augmentation for Text

- word dropout - the most simple and popular

The movie **about** cats was absolutely **great**, and the **cats** were cute.

replace with UNK

pick several words randomly


replace with random words

The movie **UNK** cats was absolutely  
**UNK**, and the **UNK** were cute.


The movie **mejorate** cats was absolutely  
**fellows**, and the **mak** were cute.

# Data Augmentation for Text

- word dropout - the most simple and popular




For images, this is similar to covering some parts: these parts are “dropped” from an image



# Data Augmentation for Text

- use external resource (e.g., thesaurus) - a bit more complicated

The **movie** about cats was **absolutely** great, and the cats were **cute**.

 pick words where you have  
synonyms and use thesaurus 

The **film** about cats was **certainly**  
great, and the cats were **nice**.

The **video** about cats was **completely**  
great, and the cats were **charming**.

# Data Augmentation for Text

- use separate models for paraphrasing - if you really care

*The movie about cats was absolutely great, and the cats were cute.*

En-Ru ↓ Yandex Translate

**Фильм о кошках был замечательный,  
и кошки были милые.**

En-Zh ↓ Google Translate

**關於貓的電影絕對很棒，  
而且貓很可愛。**

En-Ja ↓ Google Translate

**猫の映画は本当に素晴らしい、  
猫はかわいい。**

Ru-En ↓ Yandex Translate

**The cat movie was just great,  
and the cats were cute.**

Zh-En ↓ Google Translate

**Movies about cats are absolutely  
great, and cats are cute.**

Ja-En ↓ Yandex Translate

**The Cat movie is really nice and  
the cat is cute.**

# Data Augmentation for Text

- use separate models for paraphrasing - if you really care

*The movie about cats was absolutely great, and the cats were cute.*

En-Ru ↓ Yandex Translate

**Фильм о кошках был замечательный,  
и кошки были милые.**

En-Zh ↓ Google Translate

**關於貓的電影絕對很棒，  
而且貓很可愛。**

En-Ja ↓ Google Translate

**猫の映画は本当に素晴らしい、  
猫はかわいい。**

Ru-En ↓ Yandex Translate

**The cat movie was just great,  
and the cats were cute.**


Zh-En ↓ Google Translate

**Movies about cats are absolutely  
great, and cats are cute.**

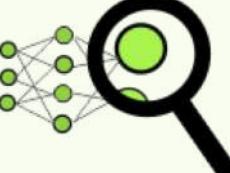
Ja-En ↓ Yandex Translate

**The Cat movie is really nice and  
the cat is cute.**


For images, this is similar to geometric transformations: we change text, but want to preserve all meaning.




# What is going to happen:

- Examples of classification tasks
- General View: Features + Classifier
- Models: Generative vs Discriminative
- Classical Methods
- Neural Methods
- Multi-Label Classification
- Practical Tips
-  Analysis and Interpretability

# What is going to happen:

- Examples of classification tasks
- General View: Features + Classifier
- Models: Generative vs Discriminative
- Classical Methods
- Neural Methods
- Multi-Label Classification
- Practical Tips
-  Analysis and Interpretability

# Analyzing Convolutional Filters




Examples of patterns captured by convolution filters for images.

The examples are from [Activation Atlas from distill.pub](#).

# Analyzing Convolutional Filters

- Find which patterns activate neurons




# Analyzing Convolutional Filters

filter	Top n-gram	Score	Top n-grams for filter 4	Score
1	poorly designed junk	7.31		
2	simply would not	5.75		
3	a minor drawback	6.11		
4	still working perfect	6.42	1 still working perfect	6.42
5	absolutely gorgeous .	5.36	2 works - perfect	5.78
6	one little hitch	5.72	3 isolation proves invaluable	5.61
7	utterly useless .	6.33	4 still near perfect	5.6
8	deserves four stars	5.56	5 still working great	5.45
9	a mediocre product	6.91	6 works as good	5.44
			7 still holding strong	5.37

A filter activates for a family of n-grams with similar meaning

The example is from the paper [Understanding Convolutional Neural Networks for Text Classification](#).

# What is going to happen:

- Examples of classification tasks
- General View: Features + Classifier
- Models: Generative vs Discriminative
- Classical Methods
- Neural Methods
- Multi-Label Classification
- Practical Tips
-  Analysis and Interpretability


# Moving Beyond Words

- Splitting text by whitespace works well in many cases.
- But this means that models have a predefined vocabulary.
- Any words not seen in the training set will be replaced by the semantically meaningless “UNK” token.
- This means we can only process a fixed number of words.

# Moving Beyond Words

- What is a good strategy when you encounter a new word though?
- Consider whether you have seen pieces of the word before!
- Remember: morphemes are the smallest units of meaning in language.
- Solution: represent rare or unknown words by their subwords

# Moving Beyond Words



# Moving Beyond Words

- This is exactly what the Byte Pair Encoding tokenization strategy does (BPE is an older algorithm, we use an adaption of it).
- Frequent words are found and kept intact while rare words are split.
- See the paper Neural Machine Translation of Rare Words with Subword Units by Rico Sennrich, Barry Haddow and Alexandra Birch

# Byte Pair Encoding

- Main Idea: Search through the current list of symbols and find the most common pairings. Then merge the pairing into a new symbol.
- There are two parts:
  - Training: find which pairs of symbols to merge
  - Inference: apply the learned merges to a new piece of text
- Let's consider each part separately.

# Byte Pair Encoding: Training

- Input: a corpus of text to train the tokenizer – not necessarily the same data we use to train our model.
- Output: a merge table and vocabulary of tokens:
  - Merge table: says which characters to combine
  - Vocabulary: which tokens we can identify – tokens are a combination of words, subwords and characters depending on the data (it's adaptive!)
- To begin we split all of the text in our corpus into their characters

# Byte Pair Encoding: Training

- With the characters split we **iteratively** apply the steps:
  - count pairs of symbols: how many times each pair occurs together in the training data;
  - find the most frequent pair of symbols;
  - merge this pair - add a merge to the merge table, and the new token to the vocabulary
- Note: we only count character pairs within words (so we can begin by counting the words to make the process more efficient)

# Byte Pair Encoding: Training - Example

- Assume we have a corpus with the following words and their number of occurrences:
  - Cat x4
  - Mat x5
  - Mats x2
  - Mate x3
  - Ate x3
  - Eat x2

# Byte Pair Encoding: Training - Example

Words (split chars)	Count
cat	4
mat	5
mats	2
mate	3
ate	3
eat	2

Merge Table
** Start Empty **

- From these counts “a” and “t” appear together the most – 19 times
- “m” and “a” appear together 10 times
- Add “at” to our merge table!

# Byte Pair Encoding: Training - Example

Words (split chars)	Count
c β	4
m β	5
m β s	2
m β e	3
β e	3
e β	2

Merge Table
at -> β

- “m” and “β” now appear together the most with 10 occurrences.
- Second is “β” and “e” with 6.
- Add “mβ” to our merge table

# Byte Pair Encoding: Training - Example

Words (split chars)	Count
c β	4
α	5
α s	2
α e	3
β e	3
e β	2

Merge Table
at -> β
mβ -> α

- Finally, we can merge “c” and “β” since this is the largest combination.

# Byte Pair Encoding: Training - Example

Words (split chars)	Count
$\mu$	4
$\alpha$	5
$\alpha s$	2
$\alpha e$	3
$\beta e$	3
$e \beta$	2

Merge Table
$at \rightarrow \beta$
$m\beta \rightarrow \alpha$
$c\beta \rightarrow \mu$

- Finally, we can merge “c” and “ $\beta$ ” since this is the largest combination.

# Byte Pair Encoding: Training - Example

Start Words (split chars)	End Words (split chars)
cat	$\mu$
mat	$\alpha$
mats	$\alpha s$
mate	$\alpha e$
ate	$\beta e$
eat	$e \beta$

- Note how the words “cat” and “mat” are common and kept together (represented with one symbol).
- Things like plurals are kept separate longer (like “mats”)
- Implementation detail: we usually add “@” to tokens which are whole words. So “mat@” and “mat s” can be treated differently.

# Byte Pair Encoding: Inference

- After learning BPE rules, you have a merge table – you can use it to segment new text.
- To start, split words into a sequence of characters. Then iteratively over the merge table and apply the merges wherever they occur in the text.
- Note that the merge table is ordered - the merges that are higher in the table were more frequent in the data and given priority.
- BPE is deterministic.