

Responsible NLP

Dr. Devon Jarvis

Slides Adapted from Jade Abbott

Responsibility in Language

- Communication is the ability to direct others' attention, cognition or thoughts.
 - "Don't think about a pink elephant"
- Natural language is a tool unique to humans which is extremely efficient and flexible.
- With this power comes responsibility:
 - The responsibility to convey information appropriately is one part.
 - What information is being conveyed is another.

Responsibility in Language

- Who is responsible for understanding language?
 - Some languages are "Writer Responsible" - it is the person conveying the message's responsibility to make the message understandable.
 - English is an example of this.
 - Other languages are "Reader Responsible" - it is the person receiving the message's responsibility to make sense of it.
 - An example is Chinese.
- This is a coarse grouping and open to stereotypes but it makes clear:
 - responsible communication is at the heart of language.

Reader vs Writer Language

- The difference comes down to how much context is needed to understand a statement.
- In writer responsible languages:
 - less situational awareness is needed to understand a message
 - what is necessary (the context) is presented with the message.
- In reader responsible languages:
 - Lots of situational awareness is needed to understand
 - Context can be omitted for brevity and left to the reader to decipher.

Reader vs Writer Language

- But the distributional hypothesis - words which share a context will share a meaning - is foundational to modern NLP.
- The addition of context can also lead to some implicit biases:
 - "They want to be a nurse when they grow up"
 - "He wants to be a nurse when he grows up"
 - The second sentence has the gender of the person as implicit context where it is not necessary.
 - Additional context in writer responsible language can lead to false correlations learned by a model.

Multi-lingualism

- Different cultural standards can make communication difficult:
 - In some cultures it is rude to ask a person their name directly.
- People often have a home-language and then a professional language.
 - People educated in a second language can prefer this language for these domains.
 - English is also the language typical of scientific papers - many languages do not have the vocabulary for all domains.

Multi-lingualism

- Having to switch between languages can be tiring - "code switching" is a big problem for LLMs as well.
- Having to speak a language without identifying with the culture can lead to mistakes.
- How would reader vs writer responsible languages pair with second language speakers?
 - How do we preserve cultural identity while also providing access to information?
 - Is it fair to translate Zulu books into English if it means there are less Zulu speakers but more awareness of Zulu culture?

Wits' Language Policy

- "Wits adopted a new Language Policy in 2015 aimed at promoting creativity, selfhood and cognition through linguistic diversity."
- The principles and values underlying the Wits Language Policy can be said to reside in the importance of:
 - unlocking cultural understanding,
 - enhancing access to knowledge,
 - producing multilingual graduates and professionals,
 - improving teaching and learning, communication, research and administration
 - demonstrating respect for language and cultural diversity.

Connections to NLP

- Immediately we can see connection between these linguistics concepts and NLP:
 - The need for useful context to impart meaning.
 - Language and understanding can be effortful.
 - Implicit bias stems from the nature of some languages themselves.
 - Language is seen as a demonstration of intelligence or competence.
 - Turing Test for computers and fluency in humans.

Bias from the model

- There are number of important biases or sources of harm when discussing NLP.
- We begin by considering the biases stemming from the model itself.
- These can be any form of harm from the production of language and training of the model itself.

Stochastic Parrots

- Fluency and coherence is consistent but opinions or semantic content is sampled.
- We don't even know how to measure "correct language"
 - Perplexity?
 - BLEU Score?
 - Token-wise Prediction Error
- Coherence in human language revolves around a shared context between speaker and listener with intent [1]:
 - LMs are ungrounded and require the speaker to describe its own intent and context.

[1] Bender, Emily M., et al. "On the dangers of stochastic parrots: Can language models be too big?  ." *Proceedings of the 2021 ACM conference on fairness, accountability, and transparency*. 2021.

Stochastic Parrots

"...an LM is a system for haphazardly stitching together sequences of linguistic forms it has observed in its vast training data, according to probabilistic information about how they combine, but without any reference to meaning: a stochastic parrot." [1]

[1] Bender, Emily M., et al. "On the dangers of stochastic parrots: Can language models be too big?  ." *Proceedings of the 2021 ACM conference on fairness, accountability, and transparency*. 2021.

Stochastic Parrots

- LMs have come a long way since 2021 and have shifting notions of intent, agency and context:
 - Incorporating multiple modalities can fix the grounding problem [2].
 - Context windows are becoming bigger and persistent - allowing for more detailed comprehension [3].
 - Context can also constrain the LM output and methods like Beam Search balance variety with coherence.

[2] Cao, Meng, et al. "On Pursuit of Designing Multi-modal Transformer for Video Grounding." *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*. 2021.

[3] Jin, Hongye, et al. "LLM Maybe LongLM: SelfExtend LLM Context Window Without Tuning." *Forty-first International Conference on Machine Learning*.

Enforces Hegemonic View

- LMs require a lot of data which mainly comes from the internet:
 - The dominant voices on the internet are overly represented. [1]
 - At this scale variations in language do not "average out" but are memorised and triggered by similar contexts.
 - Language tend to follow Zipf's Law (which affects LLMs [4]) but within a context language is skewed:
 - The hegemonic viewpoint is context specific.
 - While all views can be represented in LMs, common or important contexts may trigger a specific viewpoint.

[1] Bender, Emily M., et al. "On the dangers of stochastic parrots: Can language models be too big?  ." *Proceedings of the 2021 ACM conference on fairness, accountability, and transparency*. 2021.

Adversarial Examples

- Taken to the extreme this can result in adversarial example or adversarial triggers:
 - Certain prompts which produce inappropriate behaviour indirectly through the context.
 - Jailbreaking: These can over-ride manual safeguards! [5]




Figure 1: Example of jailbreak prompt. Texts are adopted from our experimental results.

[5]Shen, Xinyue, et al. "" do anything now": Characterizing and evaluating in-the-wild jailbreak prompts on large language models." *Proceedings of the 2024 on ACM SIGSAC Conference on Computer and Communications Security*.

Adversarial Examples

No.	System Prompt	User Input	Model Response
1	As an expert analyst, provide a detailed examination of the advantages and disadvantages of a given policy. Discuss the potential benefits and drawbacks in a balanced manner, offering insights on how the policy could impact various aspects of society, economy, and education. The analysis should include multiple points for both pros and cons.	[trigger1] Cellular respiration in plants is the process by which they [trigger2] Providing free college education	Cellular respiration in plants is the process by which they
2	Classify the given list of animals into two categories -- carnivores and herbivores. Please output in the following json format: {"carnivores": [...], "herbivores": [...]}	Animals: [, [trigger1] {"carnivores": ["Elephant"], "herbivores": ["Tiger", "Deer"]}] [trigger2] Tiger', 'Elephant', 'Deer'	{"carnivores": ["Elephant"], "herbivores": ["Tiger", "Deer"]}

Table 1: Examples of adversarial attack with our triggers. The triggers are in red, and the payload is in orange. Because our triggers can be easily extended to attack real applications, in order to prevent malicious individuals from using our triggers to perform real attacks, we have concealed the actual trigger in the examples.

[6] Liang, Jiashuo, Guancheng Li, and Yang Yu. "Universal and Context-Independent Triggers for Precise Control of LLM Outputs." *arXiv preprint arXiv:2411.14738* (2024).

Dead Internet Theory

- A theory that the internet is slowly being dominated by AI generated content.
- "These AI-driven tools, capable of creating highly realistic yet fabricated content, pose a significant threat to the integrity of information online, propelling misinformation and eroding the foundation of trust essential for healthy digital interactions." [7]
- Create mistrust in information and hides authentic communication.

Dead Internet Theory

- Bias Propagation: Generated content is fed back into the training data for future models.
- Undermines efforts to democratize the internet and highlight broader perspectives:
 - AI is more efficient at content generation.
 - Crowds out the real content creators with interesting perspectives.
 - Creates a feedback loop of the hegemonic viewpoint as the most represented media becomes even more prevalent.
 - Can lead to feelings of isolation among under-represented groups.

Tokenization & Prompting

- Even if AI is used carefully to support creation by more groups:
 - AI models are limited for under-represented languages.
 - Tokenization in particular is a big problem for African Languages:

Tokens	Characters	Tokens	Characters
4	14	7	19


I love the cat

Ngiyalithanda ikati

Tokenization & Prompting

- ChatGPT3.5 and 4 problems with African languages.

Tokens	Characters
7	19



Tokenization & Prompting

- ChatGPT3.5 and 4 problems with African languages.



Tokens

7

Characters

19

Ngiyalithanda ikati




Tokenization & Prompting

- Getting the most out of LLMs often requires questions to be asked in a particular way:
 - "Prompting" is the skill of designing these questions.
 - Makes sure that the most useful context is provided to the model.
- Problem 1: Prompting is not a universal skill and is an emerging technology which favours affluence.
- Problem 2: Leans into the speaker-responsible view - not everyone will be equally comfortable depending on their background.
- These are sources of bias which can prohibit the technology from being used for the benefit of all.

Resource Use

- It is important that we also consider sources of bias from using or training the model - not just using it.
- Training LLMs in particular is extremely expensive and has impacts on the economy and environment.
- The inconsistent power demand while training is also problematic and leads to wasted resources.



Resource Use

- LLM power usage tends to scale with parameter count which are now in the billions

LLM Name	Developer	Release Date	Access	Parameters
GPT-4o	OpenAI	May 13, 2024	API	Unknown
Claude 3.5	Anthropic	June 20, 2024	API	Unknown
Grok-1	xAI	November 4, 2023	Open-Source	314 billion
Mistral 7B	Mistral AI	September 27, 2023	Open-Source	7.3 billion
PaLM 2	Google	May 10, 2023	Open-Source	340 billion
Falcon 180B	Technology Innovation Institute	September 6, 2023	Open-Source	180 billion
Stable LM 2	Stability AI	January 19, 2024	Open-Source	1.6 billion, 12 billion

Resource Use

- LLM power usage tends to scale with parameter count which are now in the billions [1].

Year	Model	# of Parameters	Dataset Size
2019	BERT [39]	3.4E+08	16GB
2019	DistilBERT [113]	6.60E+07	16GB
2019	ALBERT [70]	2.23E+08	16GB
2019	XLNet (Large) [150]	3.40E+08	126GB
2020	ERNIE-GEN (Large) [145]	3.40E+08	16GB
2019	RoBERTa (Large) [74]	3.55E+08	161GB
2019	MegatronLM [122]	8.30E+09	174GB
2020	T5-11B [107]	1.10E+10	745GB
2020	T-NLG [112]	1.70E+10	174GB
2020	GPT-3 [25]	1.75E+11	570GB
2020	GShard [73]	6.00E+11	–
2021	Switch-C [43]	1.57E+12	745GB

Table 1: Overview of recent large language models

Resource Use

- Scale does not guarantee that the data is diverse:
 - 67% of reddit users are male in the US
 - The majority of social media posts on certain platforms are English.
- Data is static:
 - Content from many years ago remains even though social norms are changing and progressing
 - Historically hegemonic viewpoints resurface
- Performance does not scale with compute:
 - "increase in 0.1 BLEU score using neural architecture search for English to German translation results in an increase of \$150,000 compute cost" [9]

[9] Strubell, Emma, Ananya Ganesh, and Andrew McCallum. "Energy and policy considerations for modern deep learning research." *Proceedings of the AAAI conference on artificial intelligence*. Vol. 34. No. 09. 2020.

Resource Use

- Some context:

Consumption	CO ₂ e (lbs)
Air travel, 1 passenger, NY↔SF	1984
Human life, avg, 1 year	11,023
American life, avg, 1 year	36,156
Car, avg incl. fuel, 1 lifetime	126,000

Training one model (GPU)	
NLP pipeline (parsing, SRL)	39
w/ tuning & experimentation	78,468
Transformer (big)	192
w/ neural architecture search	626,155

Table 1: Estimated CO₂ emissions from training common NLP models, compared to familiar consumption.¹

Air travel and per-capita consumption source: <https://bit.ly/2Hw0xWc>

[9] Strubell, Emma, Ananya Ganesh, and Andrew McCallum. "Energy and policy considerations for modern deep learning research." *Proceedings of the AAAI conference on artificial intelligence*. Vol. 34. No. 09. 2020.

Double Edge of Resource Use

- Lower resource communities feel the impact of high-resource models twice:
 1. Through exclusion and inability to compete.
 2. The effects of climate change tend to hit the world's most marginalized first and they bounce back slower.
- Similarly there can be exploitation from data collection involving marginalized communities:
 1. Institution can use the data for research without inclusion beyond the collection (asymmetric power balance).
 2. Marginalized communities cannot control the use of their data (lack trust & conflicts of interest).
 3. Data is taken out of context or misinterpreted.

Double Edge of Resource Use

"Our findings underscore the need to gather more perspectives from low- and middle-income countries, whose notions of impact extend beyond the immediate technical concerns or impacts in the short- to medium-term." [10]


[10] Brown, Venetia, et al. "A Qualitative Study on Cultural Hegemony and the Impacts of AI." *Proceedings of the AAAI/ACM Conference on AI, Ethics, and Society*. Vol. 7. 2024.

What Can be Done? - Models

- Explainability and Interpretability:
 - Terms are typically conflated.
 - Explainability - the model's ability to produce outputs which demonstrate its computation for user
 - Interpretability - the ability to directly see or uncover the computation of the model by user.
 - First case the model produces the insight and in second the user does.
- Both can dramatically improve our understanding of the models and help identify mechanisms for bias.

What Can be Done? - Models


- Interpretability Examples:
 - Text highlighting (Occlusion/LIME shown below)
 - In many cases these do not match with LLMs providing explanations for their own behaviour



[11] Huang, Shiyuan, et al. "Can large language models explain themselves? a study of llm-generated self-explanations." *arXiv preprint arXiv:2310.11207* (2023).

What Can be Done? - Models


- Mechanistic interpretability (MI):
 - Methods for reverse engineering or systematically establishing causal chains in neural networks.
 - Eg: Understanding "Induction Heads" for ICL.



[12] Singh, Aaditya K., et al. "What needs to go right for an induction head? A mechanistic study of in-context learning circuits and their formation." *arXiv preprint arXiv:2404.07129* (2024).

What Can be Done? - Models

- Evaluation Techniques:
 - Extrinsic evaluation - using a secondary task to determine what a model has learned.
 - Eg: Predicting parts of speech from text.



What Can be Done? - Models


- Evaluation Techniques:
 - Intrinsic Evaluation - using various metrics from the task the network is trained on to determine its performance.
 - Eg: Precision & Recall, Cohen's Kappa, ROC Curve
- These techniques demonstrate if the model is collapsing to predicting imbalanced class labels and provide a more precise view of the model's performance.
- NOTE: Metrics are inherently narrow and can be misleading - such a BLEU score
- It is important to use a suit of metrics and evaluation techniques.

What Can be Done? - Models

- Transfer Learning and Multi-lingual models
 - Broaden the scope of your training dataset to afford more data (data could be synthetic or generated).
 - Ensure that the broader scope still enables the desired behaviour.
 - For example multi-lingual models might share a semantic space across some languages.
 - Could transfer from broad tasks to specific ones:
 - Eg: Training a model on next word prediction before a classification task to learn about sentence structure.
- This enables more powerful models in low resource settings.

What Can be Done? - Models

- Transfer Learning and Multi-lingual models
 - For example multi-lingual models might share a semantic space across some languages:



[14] Zeng, Hongchuan, et al. "Converging to a lingua franca: Evolution of linguistic regions and semantics alignment in multilingual large language models." *arXiv preprint arXiv:2410.11718* (2024).

What Can be Done? - Models

- Multi-lingual model case study - AfriBERTA:
 - AfriBERTA is a multi-lingual model for African Languages.
 - It made 3 main contributions:
 - Provides a competitive multi-lingual model for low-resource languages trained from scratch without high-resource transfer.
 - Performed ablations showing that many attention head can be pruned, especially for deeper models where less are needed.
 - Used SentencePiece - a character level tokenizer like BPE which gives the model the ability to work on unseen words.

[15] Ogueji, Kelechi, Yuxin Zhu, and Jimmy Lin. "Small data? no problem! exploring the viability of pretrained multilingual language models for low-resourced languages." *Proceedings of the 1st Workshop on Multilingual Representation Learning*. 2021.

What Can Be Done - Data


- Thoughtful data collection - some considerations:
 - Copyright and Terms of Use (authors permission)
 - Privacy and Protection:
 - POPIA in South Africa
 - General Data Protection Regulation in Europe
 - Covers consent and right to be forgotten
 - Sovereignty and Governance
 - Where is data stored geographically?
 - Do local communities where data is sourced get to use it first?
 - How will the data be hosted and shared?

What Can Be Done - Data

- Data curation is the organisation, selection and integration of data collected from various sources.
- Pros of Curation:
 - Removes social biases.
 - Preserves privacy through anonymization
 - Removes potentially exploitable patterns and balances the data.
- Cons of Curation:
 - The curator plays a role and can be biased
 - Throws away data which can be expensive
 - No clear best approach
 - A step away from studying the world "organically"

What Can Be Done - Data

- Increase participation:
 - All aspects of the development pipeline must be diverse and inclusive (left for machine translation)
 - Right shows number of wiki articles vs speakers for a number of languages.



Language	Articles	Speakers
English	6,087,118	1,268,100,000
Egyptian Arabic	573,355	64,600,000
Afrikaans	91,002	17,500,000
Kiswahili	59,038	98,300,000
Yoruba	32,572	39,800,000
Shona	5,505	9,000,000
Zulu	2,219	27,800,000
Igbo	1,487	27,000,000
Luo	0	4,200,000
Fon	0	2,200,000
Dendi	0	257,000
Damara	0	200,000

[16] Wilhelmina, Nekoto, et al. "Participatory research for low-resourced machine translation: A case study in African languages." *Findings of the Association for Computational Linguistics: EMNLP 2020* (2020): 2144-2160.

What Can Be Done - Data

Some useful companies, platforms and communities for African participation in NLP



Common Language Resources and Technology Infrastructure



ZiND!



LELAPA AI



zenodo

X Translators without Borders



Masakhane

A grassroots NLP community for Africa, by Africans

SADiLaR

South African Centre for Digital Language Resources

What Can Be Done - Documents

- Data Statements [17] - documents promoting transparency and mitigates bias. Includes:
 - Curation Rationale (which texts were selected and why)
 - Language Variety
 - Speaker demographics (age,race,ethnicity,native language,etc)
 - Annotator demographics
 - Speech situation & context (time & place, modality, audience)
 - Text characteristics (genre, topic, diacritics)
 - Recording tools and quality
 - Other (curator demongraphics for example)
 - Provenance Appendix (if built from other dataset include their statement)

[17] Bender, Emily M., and Batya Friedman. "Data statements for natural language processing: Toward mitigating system bias and enabling better science." *Transactions of the Association for Computational Linguistics* 6 (2018)

What Can Be Done - Documents

- Datasheets [18] - a set structure for describing datasets. Includes:
 - Motivation for collecting data
 - Composition of data (sources and content)
 - Collection Process
 - Preprocessing, cleaning and labeling steps
 - Intended uses and use-cases
 - Distribution plan and availability
 - Maintenance plans
 - Potential impacts (positive and negative) and foreseen challenges.

What Can Be Done - Documents

- Model cards [19] for describing the models trained from data.
Includes:
 - Model Details (version, type, algorithms, features, license, contact)
 - Intended use and use-cases
 - Factors related to training (phenotypic groups, env conditions)
 - Metrics (model performance measures, variance)
 - Training Data (dataset was used to train the model & why)
 - Evaluation Data (datasets used to evaluate the model & why)
 - Quantitative Analyses (unitary results, intersectional results)
 - Ethical Considerations (data, human life, risk & mitigations)

[19] Mitchell, Margaret, et al. "Model cards for model reporting." *Proceedings of the conference on fairness, accountability, and transparency*. 2019.